

Theoretical Understanding

Q1: Explain the primary differences between TensorFlow and PyTorch. When would you choose one over the other?

TensorFlow and PyTorch are two of the most widely used open-source deep learning frameworks in the machine learning community. While they serve similar purposes, each offers distinct features, workflows, and ecosystems that make them suitable for different use cases. Understanding their primary differences is crucial when deciding which to use for a particular project—whether it's cutting-edge research, rapid prototyping, or deploying models at scale.

Feature	TensorFlow	PyTorch
Developer	Google	Meta
Computation type	Static graphs	Dynamic graphs
Ease of use	Steeper learning curve	pythonic
Community support	Large and mature	Rapidly growing
Visualization	Built-in Tensor board	Tensor board among other tools
Deployment	TensorFlow Serving, TFLite, TensorFlow.js	Torch Script, ONNX
Preferred use	Production, mobile, and web deployment	Research, experimentation, fast prototyping

Choose TensorFlow when:

- You need robust tools for production, deployment, and cross-platform support.
- You're working on mobile or web apps (e.g., TFLite, [TensorFlow.js](#)).
- Visualization with TensorBoard is important.

Choose PyTorch when:

- You're doing academic research, NLP, or computer vision.
- You prefer quick debugging and a more Pythonic coding style.
- You need dynamic computation graphs for flexibility.

Q2: Describe two use cases for Jupyter Notebooks in AI development

Exploratory Data Analysis (EDA).

Jupyter Notebooks are ideal for exploring datasets, visualizing data distributions, and identifying patterns or anomalies. Data scientists use them to write code, visualize outputs (like charts and graphs), and document findings—all in one interactive environment.

Model Prototyping and Experimentation.

Jupyter allows for rapid prototyping of machine learning models. Developers can test different algorithms, adjust hyperparameters, and view real-time results, making it easier to iterate and refine AI models before moving to production.

Q3: How does spaCy enhance NLP tasks compared to basic Python string operations?

spaCy significantly enhances NLP tasks by providing **advanced, efficient, and linguistically aware tools**, which go far beyond basic Python string methods like `.split()`, `.replace()`, or `.lower()`.

The following are key enhancements:

Linguistic Awareness.

spaCy understands the structure of language. It offers features like **tokenization**, **part-of-speech tagging**, **named entity recognition (NER)**, and **dependency parsing**, which are essential for truly understanding text context—something basic string functions cannot do.

Speed and Efficiency.

spaCy is designed for industrial-strength NLP. It processes large volumes of text **faster and more accurately** than custom code written with standard Python operations.

Built-in Models and Pipelines.

spaCy comes with **pretrained models** that can recognize common entities (like names, dates, organizations), which would require substantial manual effort to handle with basic string operations.

Comparative Analysis between Scikit-learn and TensorFlow

1. **Scikit-learn** is primarily used for classical machine learning tasks such as classification, regression, clustering, and dimensionality reduction. It is highly suitable for structured (tabular) data and includes algorithms like decision trees, random forests, linear models, and support vector machines. In contrast, **TensorFlow** is designed for building and training deep learning models and is widely used for complex tasks like image recognition, natural language processing, and sequence modeling.
2. For beginners, Scikit-learn is much easier to get started with due to its simple and consistent API, making it ideal for those new to machine learning. TensorFlow, while more powerful and flexible, has a steeper learning curve—though the integration of Keras in TensorFlow 2.x has made it more accessible by simplifying model building and training.
3. In terms of community support, both libraries have strong followings. Scikit-learn has robust support within academic and traditional data science communities. TensorFlow, backed by Google, has a vast global community, frequent updates, and rich learning resources, especially for deep learning and AI research.

SCREENSHOTS

```
Command Prompt - streamlit run app.py

Requirement already satisfied: pygments<3.0.0,>=2.13.0 in c:\users\user\documents\plp academy\ai for software engineering\week 3\assignment\week_3_assignment\.venv\lib\site-packages (from rich->keras>=3.5.0->tensorflow) (2.19.1)
Requirement already satisfied: mdurl<=0.1 in c:\users\user\documents\plp academy\ai for software engineering\week 3\assignment\week_3_assignment\.venv\lib\site-packages (from markdown-it-py>=2.2.0->rich->keras>=3.5.0->tensorflow) (0.1.2)
Using cached tensorflow-2.19.0-cp311-cp311-win_amd64.whl (375.9 MB)
Using cached streamlit-1.45.1-py3-none-any.whl (9.9 MB)
Using cached altair-5.5.0-py3-none-any.whl (731 kB)
Installing collected packages: tensorflow, altair, streamlit
Successfully installed altair-5.5.0 streamlit-1.45.1 tensorflow-2.19.0

(.venv) C:\Users\user\Documents\PLP ACADEMY\AI FOR SOFTWARE ENGINEERING\WEEK 3\ASSIGNMENT\WEEK_3_ASSIGNMENT>
(.venv) C:\Users\user\Documents\PLP ACADEMY\AI FOR SOFTWARE ENGINEERING\WEEK 3\ASSIGNMENT\WEEK_3_ASSIGNMENT>streamlit run app.py

Welcome to Streamlit!

If you'd like to receive helpful onboarding emails, news, offers, promotions,
and the occasional swag, please enter your email address below. Otherwise,
leave this field blank.

Email: catherine.abugah8@gmail.com

You can find our privacy policy at https://streamlit.io/privacy-policy

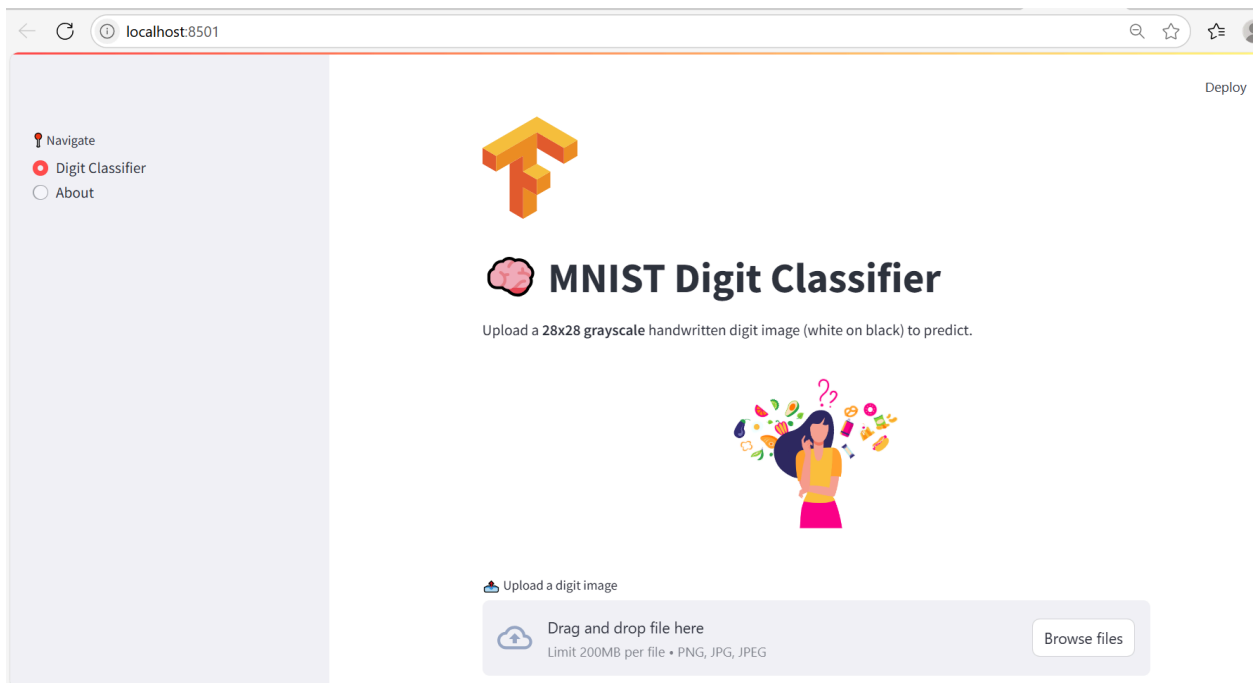
Summary:
- This open source library collects usage statistics.
- We cannot see and do not store information contained inside Streamlit apps,
  such as text, charts, images, etc.
- Telemetry data is stored in servers in the United States.
- If you'd like to opt out, add the following to %userprofile%\streamlit/config.toml,
```

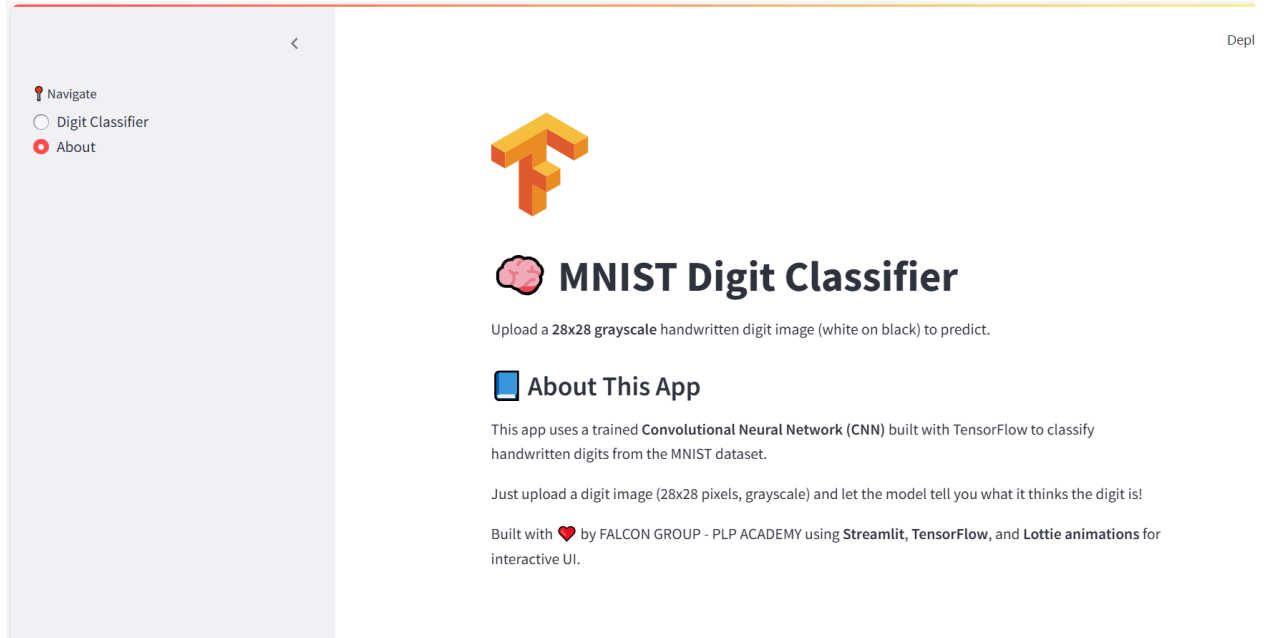
```
Command Prompt - python app.py

TensorFlow with the appropriate compiler flags.
C:\Users\user\Documents\PLP ACADEMY\AI FOR SOFTWARE ENGINEERING\WEEK 3\ASSIGNMENT\WEEK_3_ASSIGNMENT\.venv\Lib\site-packages\keras\src\saving\saving_lib.py:802: UserWarning: Skipping variable loading for optimizer 'rmsprop', because it has 1
0 variables whereas the saved optimizer has 18 variables.
  saveable.load_own_variables(weights_store.get(inner_path))
* Serving Flask app 'app'
* Debug mode: on
WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
* Running on http://127.0.0.1:5000
Press CTRL+C to quit
* Restarting with watchdog (windowsapi)
2025-06-13 05:19:56.655122: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly di
fferent numerical results due to floating-point round-off errors from different computation orders. To turn them off, se
t the environment variable 'TF_ENABLE_ONEDNN_OPTS=0'.
2025-06-13 05:19:57.478322: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly di
fferent numerical results due to floating-point round-off errors from different computation orders. To turn them off, se
t the environment variable 'TF_ENABLE_ONEDNN_OPTS=0'.
2025-06-13 05:19:58.878694: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to
use available CPU instructions in performance-critical operations.
To enable the following instructions: SSE3 SSE4.1 SSE4.2 AVX AVX2 AVX512F AVX512_VNNI FMA, in other operations, rebuild
TensorFlow with the appropriate compiler flags.
C:\Users\user\Documents\PLP ACADEMY\AI FOR SOFTWARE ENGINEERING\WEEK 3\ASSIGNMENT\WEEK_3_ASSIGNMENT\.venv\Lib\site-packages\keras\src\saving\saving_lib.py:802: UserWarning: Skipping variable loading for optimizer 'rmsprop', because it has 1
0 variables whereas the saved optimizer has 18 variables.
  saveable.load_own_variables(weights_store.get(inner_path))
* Debugger is active!
* Debugger PIN: 123-234-077
127.0.0.1 - - [13/Jun/2025 05:21:14] "GET / HTTP/1.1" 200 -
127.0.0.1 - - [13/Jun/2025 05:21:14] "GET /favicon.ico HTTP/1.1" 404 -
```

```
Command Prompt
t the environment variable `TF_ENABLE_ONEDNN_OPTS=0`.
2025-06-13 05:19:57.478322: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly di
fferent numerical results due to floating-point round-off errors from different computation orders. To turn them off, se
t the environment variable `TF_ENABLE_ONEDNN_OPTS=0`.
2025-06-13 05:19:58.878694: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to
use available CPU instructions in performance-critical operations.
To enable the following instructions: SSE3 SSE4.1 SSE4.2 AVX AVX2 AVX512F AVX512_VNNI FMA, in other operations, rebuild
TensorFlow with the appropriate compiler flags.
C:\Users\user\Documents\PLP ACADEMY\AI FOR SOFTWARE ENGINEERING\WEEK 3\ASSIGNMENT\WEEK_3_ASSIGNMENT\.venv\Lib\site-packa
ges\keras\src\saving\saving_lib.py:802: UserWarning: Skipping variable loading for optimizer 'rmsprop', because it has 1
0 variables whereas the saved optimizer has 18 variables.
  saveable.load_own_variables(weights_store.get(inner_path))
* Debugger is active!
* Debugger PIN: 123-234-077
127.0.0.1 - - [13/Jun/2025 05:21:14] "GET / HTTP/1.1" 200 -
127.0.0.1 - - [13/Jun/2025 05:21:14] "GET /favicon.ico HTTP/1.1" 404 -
1/1 _____ 0s 138ms/step
127.0.0.1 - - [13/Jun/2025 05:25:03] "POST / HTTP/1.1" 200 -
1/1 _____ 0s 31ms/step
127.0.0.1 - - [13/Jun/2025 05:28:07] "POST / HTTP/1.1" 200 -
1/1 _____ 0s 37ms/step
127.0.0.1 - - [13/Jun/2025 05:28:14] "POST / HTTP/1.1" 200 -
1/1 _____ 0s 35ms/step
127.0.0.1 - - [13/Jun/2025 05:29:04] "POST / HTTP/1.1" 200 -
```

Page





Uploaded image then here is the Prediction